

HANDLING ERRORS AND DEBUGGING

LEARNING OBJECTIVES

- Learn about how SYCL handles errors
- Learn about the difference between synchronous and asynchronous exceptions
- Learn how to handle exceptions and retrieve further information
- Learn about the host device and how to use it

SYCL EXCEPTIONS

- In SYCL errors are handled by throwing exceptions.
- It is crucial that these errors are handled, otherwise your application could fail in unpredictable ways.
- In SYCL there are two kinds of error:
 - Synchronous errors (thrown in user thread) .
 - Asynchronous errors (thrown by the SYCL scheduler).

HANDLING ERRORS

```
int main() {  
    queue q();  
  
    /* Synchronous code */  
  
    q.parallel_for(buf0.get_range(), [=](id<1> i) {  
        /* Asynchronous code */  
    });  
});  
}
```

- Kernels run asynchronously on the device, and will throw asynchronous errors.
- Everything else runs synchronously on the host, and will throw synchronous errors.

SYCL EXCEPTIONS

Synchronous
exceptions

SYCL interface

Asynchronous
exceptions

SYCL Runtime

(optional)
CPU device

Kernel
loader

Runtime
Scheduler

Data dependency
tracker

Backend interface (e.g. OpenCL API)

HANDLING ERRORS

```
int main() {  
    queue q();  
  
    /* Synchronous code */  
  
    q.submit([&](handler &cgh) {  
        /* Synchronous code */  
  
        cgh.single_task([=](id<1> i) {  
            /* Asynchronous code */  
        });  
    }).wait();  
}
```

- Code on the device runs asynchronously
- If errors are not handled, the application can fail:
 - SYCL 1.2.1 application will fail silently.
 - SYCL 2020 provides a default async handler that will call `std::terminate` when an asynchronous error is thrown.

```
1  int main() {
2      std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0,
3      0, 0, 0 };
4      try {
5          queue gpuQueue(gpu_selector{});
6
7          buffer bufA{dA};
8          buffer bufB{dB};
9          buffer bufO{dO};
10
11         gpuQueue.submit([&](handler &cgh) {
12             auto inA = accessor{bufA, cgh, read_only};
13             auto inB = accessor{bufB, cgh, read_only};
14             auto out = accessor{bufO, cgh, write_only};
15
16             cgh.parallel_for(bufO.get_range(), [=](id<1> i) {
17                 out[i] = inA[i] + inB[i];
18             });
19         }).wait();
20     } catch (...) { /* handle errors */ }
21 }
```

- Synchronous errors are typically thrown by SYCL API functions.

to handle all SYCL errors you must wrap everything in a try-catch block.

```
1  int main() {
2      std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0,
3      0, 0, 0 };
4      try{
5          queue gpuQueue(gpu_selector{}, async_handler{});
6          buffer bufA{dA};
7          buffer bufB{dB};
8          buffer bufO{dO};
9
10         gpuQueue.submit([&](handler &cgh) {
11             auto inA = accessor{bufA, cgh, read_only};
12             auto inB = accessor{bufB, cgh, read_only};
13             auto out = accessor{bufO, cgh, write_only};
14
15             cgh.parallel_for(bufO.get_range(), [=](id<1> i) {
16                 out[i] = inA[i] + inB[i];
17             });
18         }).wait();
19
20         gpuQueue.throw_asynchronous();
21     } catch (...) { /* handle errors */
22     }
```

- Asynchronous errors that may have occurred will be thrown after a command group has been submitted to a queue.


```
1  int main() {
2      std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0,
3      0, 0, 0 };
4      try{
5          queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
6              for (auto e : eL) { std::rethrow_exception(e); }
7          });
8          buffer bufA{dA};
9          buffer bufB{dB};
10         buffer bufO{dO};
11
12         gpuQueue.submit([&](handler &cgh) {
13             auto inA = accessor{bufA, cgh, read_only};
14             auto inB = accessor{bufB, cgh, read_only};
15             auto out = accessor{bufO, cgh, write_only};
16
17             cgh.parallel_for(bufO.get_range(), [=](id<1> i) {
18                 out[i] = inA[i] + inB[i];
19             });
20         }).wait();
21
22         gpuQueue.throw_asynchronous();
23     } catch (...) { /* handle errors */ }
24 }
```

```
1  int main() {
2      std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, d0{ 0,
3      0, 0, 0 };
4      try {
5          queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
6              for (auto e : eL) { std::rethrow_exception(e); }
7          });
8          ...
9
10         gpuQueue.throw_asynchronous();
11     } catch (const std::exception& e) {
12         std::cout << "Exception caught: " << e.what()
13         << std::endl;
14     }
15 }
```

- Once rethrown and caught, a SYCL exception can provide information about the error
- The `what` member function will return a string with more details

EXCEPTION TYPES

- In SYCL 1.2.1 there are a number of different exception types that inherit from `std::exception`
 - E.g. `runtime_error`, `kernel_error`
- SYCL 2020 only has a single `sycl::exception` type which provides different error codes
 - E.g. `errc::runtime`, `errc::kernel`

DEBUGGING SYCL KERNEL FUNCTIONS

- SYCL requires: a device must always be available
- → Implementations provide a CPU device
 - In AdaptiveCpp, this is the OpenMP backend
 - ⇒ Debug SYCL kernels: use the CPU device + a standard C++ debugger (gdb/lldb)

- In general, users can query the `host_debuggable` device aspect to check whether they can use the same functionality

```
1  int main() {
2      std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0,
3      0, 0, 0 };
4      try{
5          queue hostQueue(aspect_selector<aspect::host_debuggable>(),
6          async_handler{});
7
8          buffer bufA{dA};
9          buffer bufB{dB};
10         buffer bufO{dO};
11
12         hostQueue.submit([&](handler &cgh) {
13             auto inA = accessor{bufA, cgh, read_only};
14             auto inB = accessor{bufB, cgh, read_only};
15             auto out = accessor{bufO, cgh, write_only};
16
17             cgh.parallel_for(bufO.get_range(), [=](id<1> i) {
18                 out[i] = inA[i] + inB[i];
19             });
20         });
21     } catch (...) { /* handle errors */ }
```


QUESTIONS

EXERCISE

Code_Exercises/Exercise_4_Section_4_Handling_Errors/source

Add error handling to a SYCL application for both synchronous and asynchronous errors.

